
Customer Segmentation Platform Analysis Sauti.

BY: KYLE KEHOE, ANOUSHKA NAYAK, PRIYANSHA RASTOGI,
JUSTIN CHOW

ANALYTICS TEAM



KYLE

OMSBA
Spring 2023



ANOUSHKA

OMSBA
Spring 2023



PRIYANSHA

OMSBA
Spring 2023



JUSTIN

OMSBA
Fall 2022

TABLE OF CONTENTS

01 Objective

02 Data Architecture

03 Exploratory Data
Analysis

04 Consumer Behavior
Forecast

05 Demographic Clustering

06 Problems Encountered

07



01

Objective

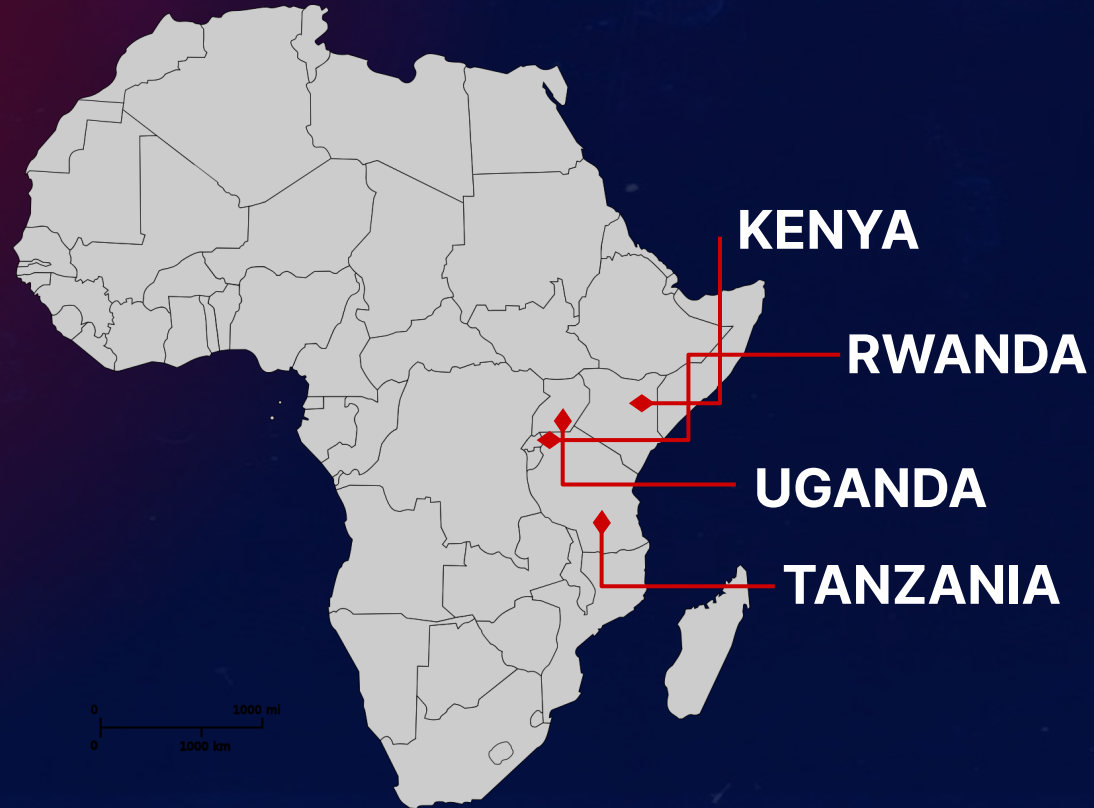
TEAM OBJECTIVE



Better understanding the analytical context of Sauti, our team objective is to better describe and understand the past and current customer behavior across their user base. Key demographics include the following:

- Gender
- Age
- Occupation
- Education
- Etc.

SAUTI EAST AFRICA REGIONS



The background is a dark blue gradient. A horizontal red line is positioned near the top. A vertical red line is on the right side. Another vertical red line is on the left side, partially obscured by a dark red circle. The text '02' is in large white font.

02

Data Architecture

DATA PACKAGE CONTENTS I

PLATFORM_SESSIONS

- **sess_id** : Unique field that tracks user sessions
- **cell_num_id** : User phone number
- **created_date** : the date the session was initialized
- **update** : the last session interaction
- **data** : session data collected in php serialized format
- **platform_id** : info platform accessed by the user.

PLATFORM_REQUESTS

- **request_id** : Unique ID
- **request_text** : user submitted input to Sauti
- **update** : time the record was submitted
- **sess_id** : session request is associated with

DATA PACKAGE CONTENTS II

PLATFORM_REPLIES

- **reply_id**: unique ID
- **request_id**: associated request for which this record is a reply
- **response_id**: record for which the content of the reply
- **udate**: time this record was submitted
- **sess_id**: session this request is associated with

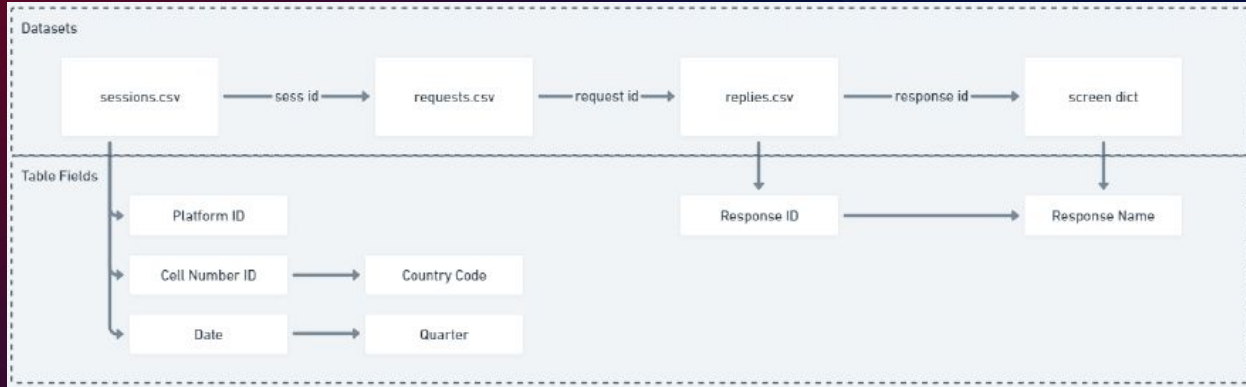
PLATFORM_SCREEN_DICT

- **platform_id**: platform that this response is associated with
- **response_id**
- **response_theme**: theme of our users' response
- **level**: organizational level of this response
- **parent**: parent to this response or top level

PLATFORM_DATABANK

- **cell_num_id**: user phone number
- **sess_id**: session ID from platform sessions where data is recorded from the user inputs
- **key_name**: collection of user-based and session-based key values (e.g., gender, education, crossingfreq, occupation, etc.)
- **value_name**: translated value
- **created_date**: time the value was recorded

DATA ARCHITECTURE WORKFLOW



- **session.csv:** if USER does not have an active Session ID, create new Session ID in table ``platform_sessions``
- **request.csv:** Record what the user sent to Sauti in the table ``platform_requests``
- **replies.csv:** Calculates the appropriate response to send to user from the table ``platform_responses``. Insert into ``platform_replies`` table.
- **screen_dict.csv:** Response IDs are thematically categorized in the table ``platform_screen_dict``.

Data collected in ``platform_session_data`` is processed, cleaned, and saved into ``platform_databank``.

The background is a dark blue gradient. A large, semi-transparent dark red circle is positioned in the bottom right corner. Three thin red lines are present: a horizontal line near the top center, a vertical line on the left side, and a vertical line on the right side.

03

Exploratory Data Analysis

EXPLORATORY DATA ANALYSIS

Complete the definition of the following concepts

STAGE 1

DATA COLLECTION

Collect platform
user and session
data from Sauti

STAGE 2

DATA CLEANING

Cleaning data to
gather exact data
required for analysis

STAGE 3

ANALYZE

Analyze and explore
aggregate data for
conclusive findings

STAGE 4

VISUALIZE RESULTS

Visualize results in
graphic form (e.g.,
bar, line, pie chart)

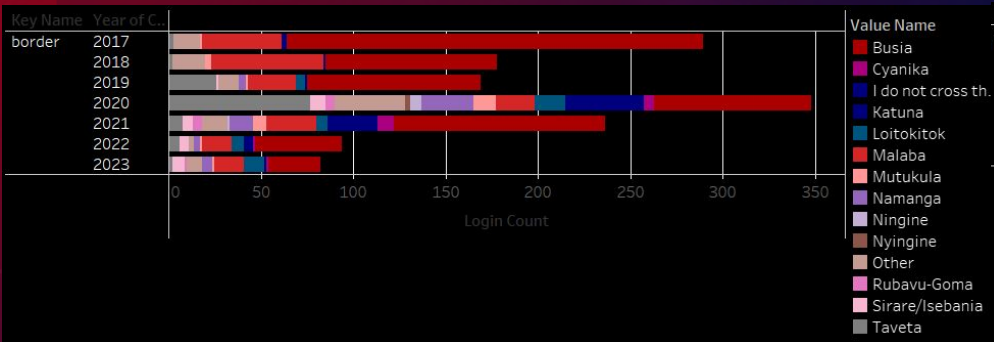
User vs Session Key Analysis

User-based analysis is distribution of the users attributes, while session-based analysis is the user participating in a session per year.



USER-BASED ANALYSIS

USER SESSIONS VS BORDER



USER SESSIONS VS GENDER (M/F)



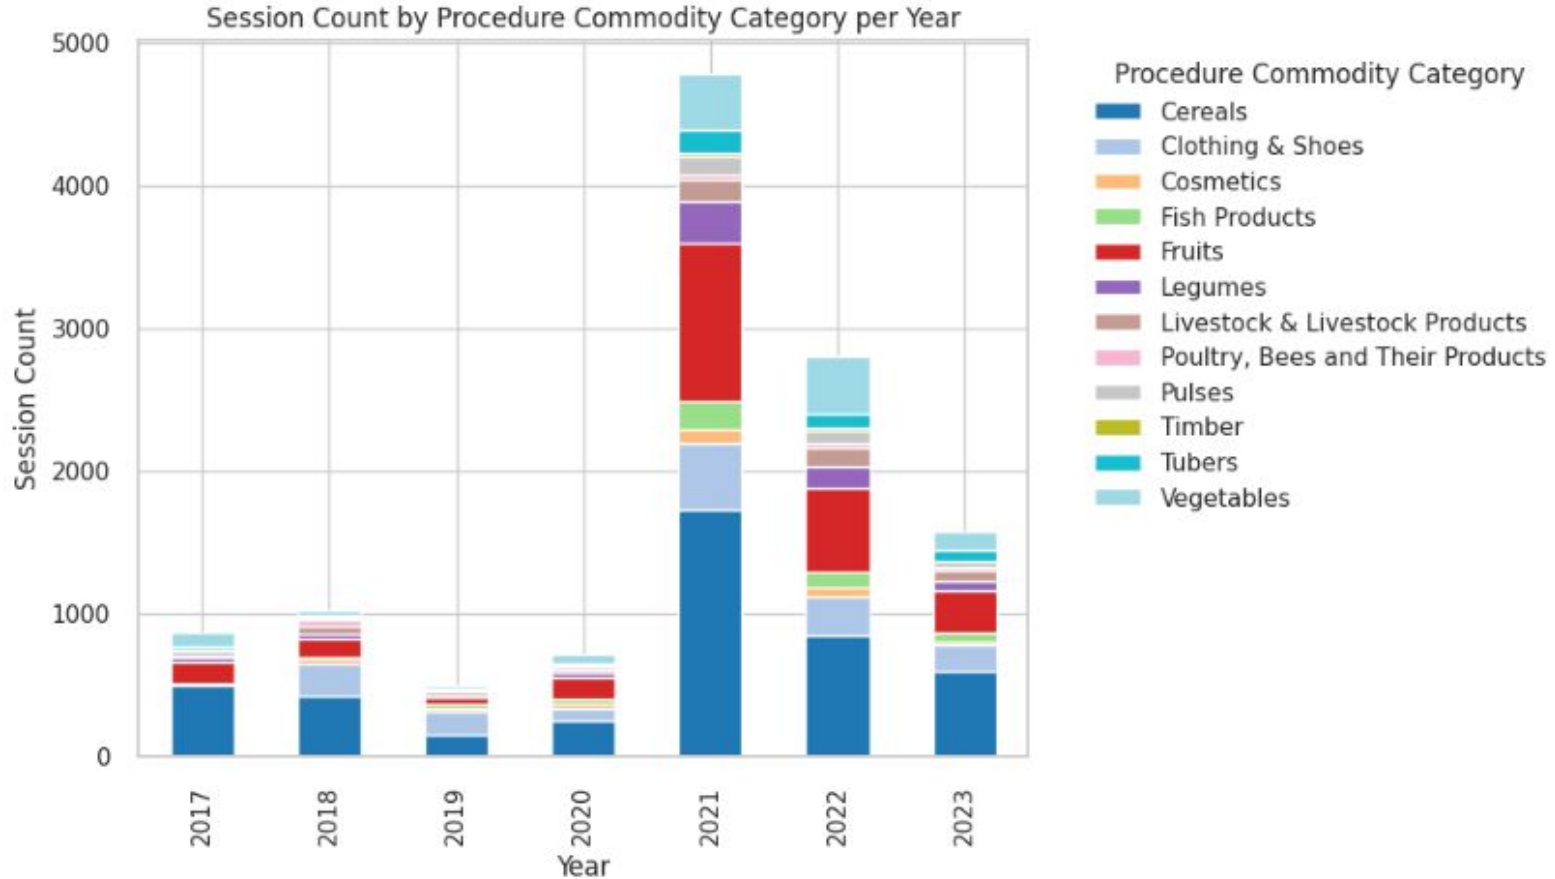
USER SESSION VS AGE RANGE



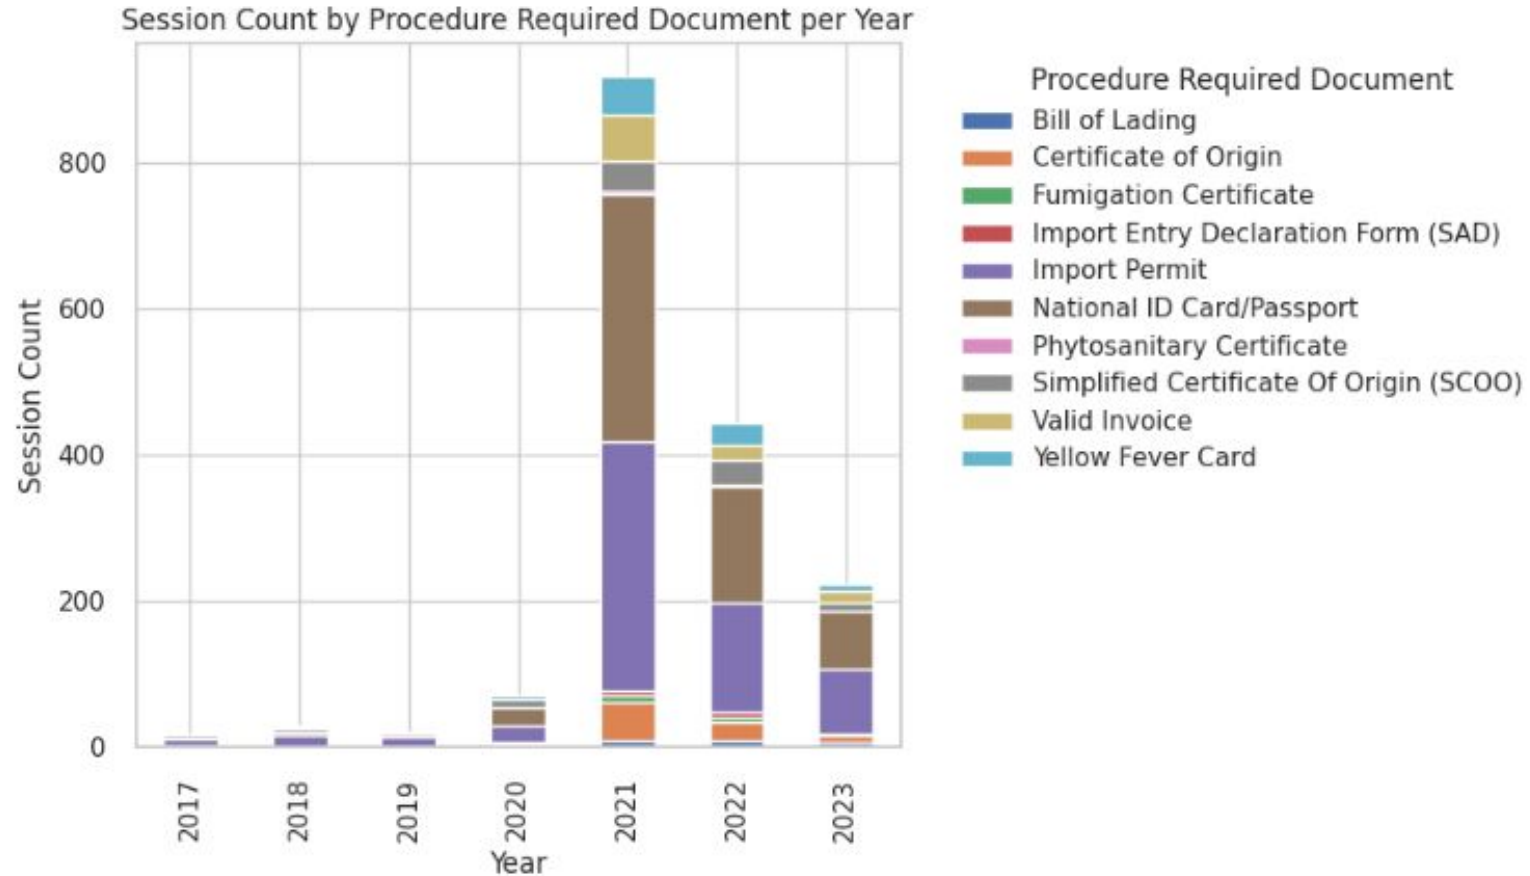
USER SESSION VS EDUCATION LEVEL




SESSION-BASED ANALYSIS



SESSION-BASED ANALYSIS




SESSION-BASED ANALYSIS




SOCIAL DETERMINANT ANALYSIS

TOTAL SESSIONS VS AGE GROUP

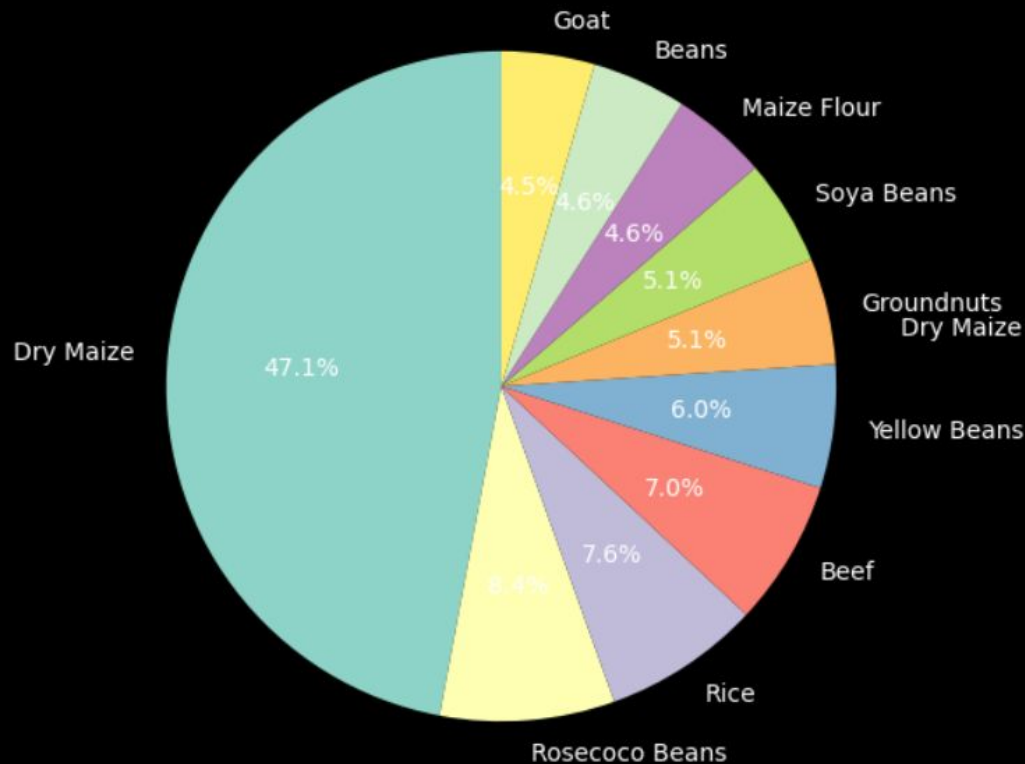


TOTAL SESSIONS VS COMMODITY CATEGORY

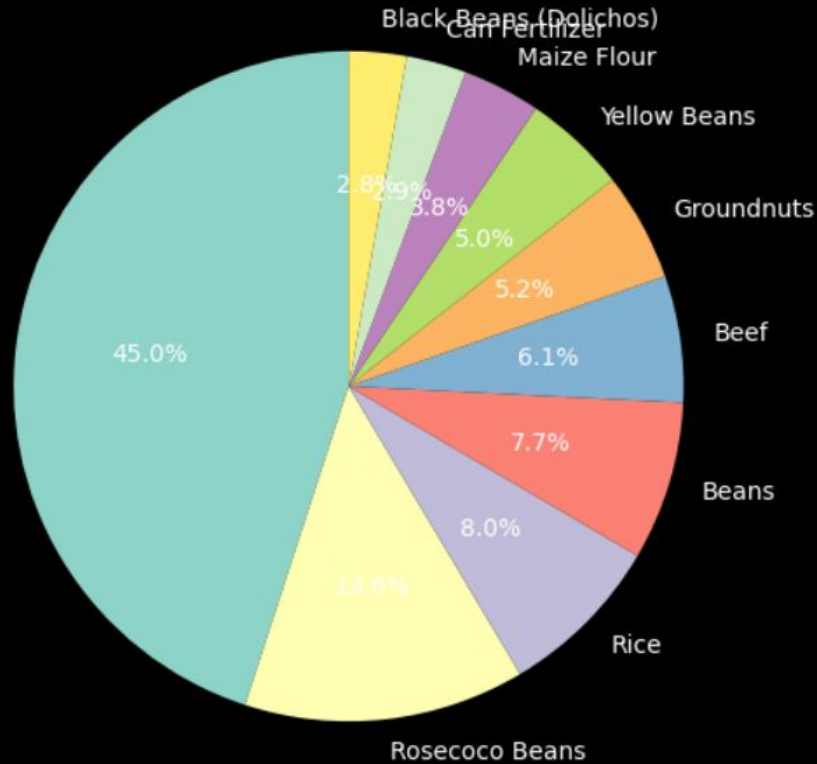


SOCIAL DETERMINANT ANALYSIS

Top 10 Commodity Products for Male




Top 10 Commodity Products for Female

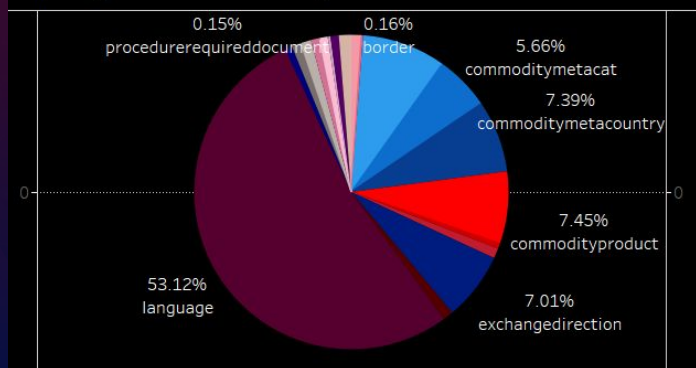


COMPREHENSIVE DASHBOARD

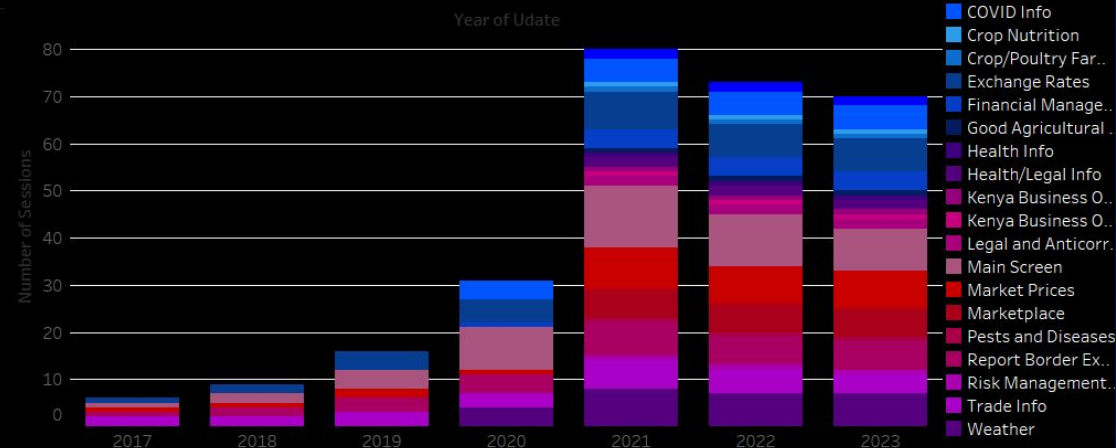
Top Response Themes Percentages




Top Key Name Percentages



Response Theme vs Num of Sessions (Per Year)



Number of Entries Over Time (Monthly Basis)



04

Consumer Behavior Forecast: Commodity Meta Categories

Probability Prediction Algorithm for Commodity
Meta Category using Random Forest Classifier
ML model.

PRE-PROCESSING DATA

Step 1:

Create Pivot Tables to reshape the data based on specific columns, creating a new Dataframe with index as Sess_id, the key name column as columns and value_name column as the values.

Step 3:

Drop Sess_id column, not needed for analysis for this part.

Step 2:

Remove duplicate entries based on sess_id and key_name, using the agg function, keeping the first non null value

Step 4:

Remove null value rows for gender , age and commodity meta category columns and one hot encode gender and age columns

PROBABILITY DISTRIBUTION

```
# Filter relevant rows and pivot the data
db_rf = databank[databank['key_name'].isin(['gender', 'age', 'commoditymetacat'])]

# Remove duplicate entries based on 'sess_id'
db_rf = db_rf.drop_duplicates(subset=['sess_id', 'key_name'], keep='first')

# databank = databank.pivot_table(index='sess_id', columns='key_name', values='value_name', aggfunc='first')
db_rf = db_rf.pivot(index='sess_id', columns='key_name', values='value_name')
db_rf.reset_index(inplace=True)

db_rf = db_rf.drop(['sess_id', axis=1])

db_rf = db_rf[pd.notna(db_rf['gender'])]
db_rf = db_rf[pd.notna(db_rf['age'])]
db_rf = db_rf[pd.notna(db_rf['commoditymetacat'])]

db_rf = pd.get_dummies(db_rf, columns=['gender', 'age'], prefix=['gender', 'age'])

X = db_rf.drop('commoditymetacat', axis=1)
y = db_rf['commoditymetacat']
```



```
[ ] # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train a Random Forest classifier (you can choose other classifiers)
model = RandomForestClassifier(n_estimators=100)
model.fit(X_train, y_train)
```

```
] # Predict Probability Distribution
# Input age range and gender for prediction
input_data = X_train.columns.to_list() # Use the same columns that were used for training
print('Training columns:' + str(input_data))

input_data = pd.DataFrame({
    'gender_Female': [1],
    'gender_Male': [0],
    'age_21-30': [1],
    'age_31-40': [0],
    'age_41-50': [0],
    'age_51-60': [0],
    'age_61-70': [0],
    'age_<21': [0],
})

# Predict probability distribution of commoditymetacat
probs = model.predict_proba(input_data)

# Get the class labels (commoditymetacat names)
class_labels = model.classes_
```

PREDICTION RESULTS

```
# Print the names and probabilities
for label, prob in zip(class_labels, probs[0]):
    print(f"Commodity: {label}, Probability: {prob:.2f}")
```

```
Training columns:['gender_Female', 'gender_Male', 'age_21-30', 'age_31-40', 'age_41-50', 'age_51-60', 'age_61-70', 'age_<21']
Commodity: Cereals, Probability: 0.41
Commodity: Farm Inputs, Probability: 0.02
Commodity: Fruits & Nuts, Probability: 0.15
Commodity: Livestock & Animal Products, Probability: 0.21
Commodity: Other, Probability: 0.00
Commodity: Other Food Products, Probability: 0.02
Commodity: Vegetables & Tubers, Probability: 0.20
```

05

Demographic Clustering for Customer Segmentation

Preprocess the data with an LLM, run the embeddings through K-means, finally use Light GBM to explain each cluster

PREPROCESSING DATA

1. Split Key/Value Pairs

- Pivot tables suggests that the data may need to be organized in a tabular format
- The process involves splitting key-value pairs
- Restructure or normalize the data.

2. Convert Text to Embedded

- Reference to the process of representing text data as numerical vectors
- Common natural al language processing tasks such as sentiment analysis, text classification, and language modeling require this step.

PREPROCESSING DATA

```
[39] model = SentenceTransformer(r"sentence-transformers/paraphrase-MiniLM-L6-v2")
      output = model.encode(sentences=text,
                             show_progress_bar=True,
                             normalize_embeddings=True)

      df_embedding = pd.DataFrame(output)
      df_embedding
```


Batches: 100%  11544/11544 [01:58<00:00, 90.75it/s]

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | ... | 374 | 375 | 376 | 377 | 378 | 379 | 380 | 381 | |
|--------|----------|----------|-----------|-----------|-----------|-----------|----------|-----------|-----------|----------|-----|----------|-----------|-----------|-----------|-----------|----------|----------|-----------|--------|
| 0 | 0.054630 | 0.031877 | 0.000613 | -0.029122 | -0.063190 | -0.032790 | 0.026601 | -0.007920 | -0.022709 | 0.062642 | ... | 0.053489 | -0.065470 | -0.002234 | -0.042182 | -0.034440 | 0.010583 | 0.050081 | -0.008227 | -0.015 |
| 1 | 0.058280 | 0.028668 | -0.002655 | -0.034798 | -0.054955 | -0.033341 | 0.029350 | -0.001681 | -0.028377 | 0.065665 | ... | 0.058561 | -0.070146 | -0.003766 | -0.024526 | -0.030072 | 0.010911 | 0.039144 | -0.003534 | -0.006 |
| 2 | 0.058280 | 0.028668 | -0.002655 | -0.034798 | -0.054955 | -0.033341 | 0.029350 | -0.001681 | -0.028377 | 0.065665 | ... | 0.058561 | -0.070146 | -0.003766 | -0.024526 | -0.030072 | 0.010911 | 0.039144 | -0.003534 | -0.006 |
| 3 | 0.054630 | 0.031877 | 0.000613 | -0.029122 | -0.063190 | -0.032790 | 0.026601 | -0.007920 | -0.022709 | 0.062642 | ... | 0.053489 | -0.065470 | -0.002234 | -0.042182 | -0.034440 | 0.010583 | 0.050081 | -0.008227 | -0.015 |
| 4 | 0.058280 | 0.028668 | -0.002655 | -0.034798 | -0.054955 | -0.033341 | 0.029350 | -0.001681 | -0.028377 | 0.065665 | ... | 0.058561 | -0.070146 | -0.003766 | -0.024526 | -0.030072 | 0.010911 | 0.039144 | -0.003534 | -0.006 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 369373 | 0.055428 | 0.027773 | -0.003660 | -0.037542 | -0.053768 | -0.035492 | 0.028892 | -0.002332 | -0.030077 | 0.068247 | ... | 0.059260 | -0.071159 | -0.004889 | -0.027548 | -0.030264 | 0.012900 | 0.040368 | -0.007033 | -0.004 |
| 369374 | 0.047667 | 0.031672 | -0.014526 | -0.036547 | -0.054998 | -0.031646 | 0.025585 | -0.007926 | -0.028124 | 0.069366 | ... | 0.053428 | -0.071963 | 0.002693 | -0.037297 | -0.031287 | 0.024995 | 0.046671 | -0.007309 | -0.011 |
| 369375 | 0.055428 | 0.027773 | -0.003660 | -0.037542 | -0.053768 | -0.035492 | 0.028892 | -0.002332 | -0.030077 | 0.068247 | ... | 0.059260 | -0.071159 | -0.004889 | -0.027548 | -0.030264 | 0.012900 | 0.040368 | -0.007033 | -0.004 |
| 369376 | 0.007176 | 0.044907 | -0.027020 | -0.029548 | -0.058843 | -0.042840 | 0.021809 | -0.002338 | -0.019830 | 0.071603 | ... | 0.064967 | -0.064455 | 0.020173 | -0.035466 | -0.035795 | 0.009136 | 0.051855 | -0.035575 | -0.019 |
| 369377 | 0.055428 | 0.027773 | -0.003660 | -0.037542 | -0.053768 | -0.035492 | 0.028892 | -0.002332 | -0.030077 | 0.068247 | ... | 0.059260 | -0.071159 | -0.004889 | -0.027548 | -0.030264 | 0.012900 | 0.040368 | -0.007033 | -0.004 |

369378 rows x 384 columns

K-MEANS CLUSTERING


```
[40] km = KMeans(init="k-means++", random_state=0, n_init="auto")  
visualizer = KElbowVisualizer(km, k=(2,20), locate_elbow=False)  
  
visualizer.fit(df_embedding)  
visualizer.show()
```



```
[45] print(f"Silhouette Score: {silhouette_score(df_embedding, clusters_predict)}")
```

➡ Silhouette Score: 0.58840411901474

Utilize Light GBM to show the most important variables for determining cluster, and predict the “average” profile in each one



INTERPRETATION (CONT.)

| key_name | cluster | age | border | occupation | gender | education | crossingfreq | produce | commodityproduct | commoditymarket | language | proceduredest |
|----------|---------|-------|---------------------------|------------------------------------|--------|-----------|--------------|---------|------------------|-----------------|----------|---------------|
| 0 | 1 | 21-30 | Busia | Trader | Male | Secondary | Never | Yes | Beef | Busia | English | UGA->KEN |
| 1 | 3 | 21-30 | Busia | Farming, Fishing, Animal Husbandry | Male | Secondary | Never | Yes | Dry Maize | Busia | English | UGA->KEN |
| 2 | 7 | 21-30 | Busia | Trader | Male | Primary | Daily | No | Beef | Nairobi | Somali | KEN->TZA |
| 3 | 4 | 31-40 | Busia | Trader | Male | Secondary | Never | Yes | Dry Maize | Busia | English | UGA->KEN |
| 4 | 0 | <21 | Busia | Other | Female | Primary | Never | Yes | Rice | Kampala | Swahili | UGA->KEN |
| 5 | 2 | <21 | Busia | Farming, Fishing, Animal Husbandry | Male | Secondary | Never | Yes | Rosecoco Beans | Busia | English | UGA->KEN |
| 6 | 5 | <21 | Busia | Trader | Male | Secondary | Never | Yes | Rice | Lira | English | UGA->KEN |
| 7 | 6 | <21 | I do not cross the border | Other | Male | Primary | Never | Yes | Rice | Mbeya | Swahili | KEN->TZA |

The background is a dark blue gradient. A large, semi-transparent dark red circle is positioned in the bottom right corner. Three thin red lines are present: a horizontal line near the top left, and two vertical lines, one on the left and one on the right, both extending from the bottom towards the middle of the slide.

06

Problems Encountered

PROBLEM: MISSING VALUES

| | key_name | age | border | oomcommoditymarket | oomcommoditymetaoat | oomcommoditymetaocountry | oomcommodityproduct | crossingfreq | eduocation | exohangedireotlon | gender | language | o |
|---------|--------------|-----|--------|--------------------|---------------------|--------------------------|---------------------|--------------|------------|-------------------|--------|----------|---|
| sess_id | oell_num_id | | | | | | | | | | | | |
| 50.0 | 254000000003 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 52.0 | 254000000012 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 54.0 | 254000000010 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 59.0 | 254000000007 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 61.0 | 254000000004 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |

| crossingfreq | education | ... | gender | language | occupation | procedurecommodity | procedurecommoditycat | proceduredest | procedureorigin | procedurerelevantagency | procedurerequireddocument | prod |
|--------------|-----------|-----|--------|----------|------------|--------------------|-----------------------|---------------|-----------------|-------------------------|---------------------------|------|
| NaN | NaN | ... | NaN | NaN | NaN | Maize | Cereals | UGA->KEN | EAC | NaN | NaN | |
| NaN | NaN | ... | NaN | NaN | NaN | Maize | Cereals | NaN | NaN | NaN | NaN | |
| NaN | NaN | ... | NaN | NaN | NaN | Maize | Cereals | NaN | NaN | NaN | NaN | |
| NaN | NaN | ... | NaN | NaN | NaN | Maize | Cereals | UGA->KEN | EAC | NaN | NaN | |
| NaN | NaN | ... | NaN | NaN | NaN | Maize | Cereals | NaN | NaN | NaN | NaN | |

PROBLEM EXPLANATION

Data Inconsistency

There is inconsistency in the data; for example, the highlighted cells under `procedurecommodity` and `procedurecommoditycat` show "Maize" and "Cereals" which might indicate a need for standardization if they are meant to represent similar categories.

Lack of Identifier Uniqueness

If the `sess_id` or `oell_num_id` columns are meant to be unique identifiers, the presence of NaN could be problematic for tracking individual sessions or items.

Categorization Issue

If 'Maize' is a specific type of 'Cereal', then it may be correctly categorized under a broader 'Cereals' category. However, if there are other specific commodities listed that should also fall under 'Cereals' but do not, this could indicate inconsistent categorization.

PROBLEM INVESTIGATION

| | A | B | C | D | E | F | G |
|------|--------------|---------|-----------------------|----------------|--------------|---|---|
| 1 | cell num id | sess id | key_name | value_name | created_date | | |
| 14 | 254000000004 | 61 | procedurecommoditycat | Cereals | 48:22.0 | | |
| 15 | 254000000004 | 61 | procedurecommodity | Maize | 48:22.0 | | |
| 172 | 254000000004 | 217 | commoditymarket | Eldoret | 12:40.0 | | |
| 173 | 254000000004 | 217 | commodityproduct | Dry Maize | 12:40.0 | | |
| 174 | 254000000004 | 217 | procedurecommoditycat | Cereals | 12:40.0 | | |
| 175 | 254000000004 | 217 | procedurecommodity | Rice - Husked | 12:40.0 | | |
| 176 | 254000000004 | 217 | proceduredest | UGA->KEN | 12:40.0 | | |
| 177 | 254000000004 | 217 | procedureorigin | EAC | 12:40.0 | | |
| 823 | 254000000004 | 791 | commoditymarket | Dodoma | 15:20.0 | | |
| 824 | 254000000004 | 791 | commodityproduct | Sunflower Seed | 15:20.0 | | |
| 863 | 254000000004 | 825 | procedurecommoditycat | Cereals | 07:58.0 | | |
| 864 | 254000000004 | 825 | procedurecommodity | Maize | 07:58.0 | | |
| 865 | 254000000004 | 825 | proceduredest | UGA->KEN | 07:58.0 | | |
| 866 | 254000000004 | 825 | procedureorigin | EAC | 07:58.0 | | |
| 1267 | 254000000004 | 1113 | commoditymarket | Eldoret | 15:45.0 | | |
| 1268 | 254000000004 | 1113 | commodityproduct | Dry Maize | 15:45.0 | | |
| 1473 | 254000000004 | 1339 | commoditymarket | Kitale | 45:14.0 | | |
| 1474 | 254000000004 | 1339 | commodityproduct | Dry Maize | 45:14.0 | | |
| 1485 | 254000000004 | 1347 | procedurecommoditycat | Cereals | 16:08.0 | | |
| 1486 | 254000000004 | 1347 | procedurecommoditycat | Cereals | 16:08.0 | | |
| 1487 | 254000000004 | 1347 | procedurecommoditycat | Cereals | 16:08.0 | | |
| 1488 | 254000000004 | 1347 | procedurecommodity | Maize | 16:08.0 | | |
| 1489 | 254000000004 | 1347 | procedurecommodity | Maize | 16:08.0 | | |
| 1490 | 254000000004 | 1347 | proceduredest | KEN->UGA | 16:08.0 | | |
| 1491 | 254000000004 | 1347 | proceduredest | UGA->KEN | 16:08.0 | | |
| 1492 | 254000000004 | 1347 | procedureorigin | EAC | 16:08.0 | | |
| 1493 | 254000000004 | 1347 | procedureorigin | EAC | 16:08.0 | | |
| 1524 | 254000000004 | 1380 | commoditymarket | Owino | 02:00.0 | | |
| 2420 | 254000000004 | 2185 | language | English | 09:26.0 | | |
| 2421 | 254000000004 | 2185 | commoditymarket | Kitale | 09:26.0 | | |
| 2422 | 254000000004 | 2185 | commodityproduct | Dry Maize | 09:26.0 | | |
| 2535 | 254000000004 | 2232 | language | English | 42:20.0 | | |
| 2536 | 254000000004 | 2232 | exchangedirection | KES->UGX | 42:20.0 | | |

INVESTIGATION EXPLANATION

Hierarchy and Classification

'Cereals' appears in the column labeled 'procedurecommoditycat', suggesting it's a category for commodities. 'Maize' appears under 'procedurecommodity', indicating that it is a specific commodity within that category.

Data Duplication

The repeated appearance of '254000000004' under 'cell_num_id' with different 'procedurecommodity' and 'procedurecommoditycat' entries could indicate that the dataset contains multiple transactions or records for the same entity. This is explained by multiple users with different `sess_id` using the same cell phone.

PROBLEM COMPARISON

| | A | B | C | D | E | F | | | | | | | | | | | | | | | | |
|----|--------------|---------|-----------------------|----------------|--------------|---|----------|----------|---------------|-----|--------|-----------------|------------------|----------------------|------------------|--------------|-----------|-----|--------|----------|------------|-----------|
| | cell_num_id | sess_id | key_name | value_name | created_date | | | | | | | | | | | | | | | | | |
| | 254000000004 | 61 | procedurecommoditycat | Cereals | 48:22.0 | | | | | | | | | | | | | | | | | |
| | 254000000004 | 61 | procedurecommodity | Maize | 48:22.0 | | | | | | | | | | | | | | | | | |
| 2 | 254000000004 | 217 | commoditymarket | Eldoret | 12:40.0 | | | | | | | | | | | | | | | | | |
| 3 | 254000000004 | 217 | commodityproduct | Dry Maize | 12:40.0 | | | | | | | | | | | | | | | | | |
| 4 | 254000000004 | 217 | procedurecommoditycat | Cereals | 12:40.0 | | | | | | | | | | | | | | | | | |
| 5 | 254000000004 | 217 | procedurecommodity | Rice - Husked | 12:40.0 | | | | | | | | | | | | | | | | | |
| 6 | 254000000004 | 217 | proceduredest | UGA->KEN | 12:40.0 | | | | | | | | | | | | | | | | | |
| 7 | 254000000004 | 217 | procedureorigin | EAC | 12:40.0 | | | | | | | | | | | | | | | | | |
| 8 | 254000000004 | 791 | commoditymarket | Dodoma | 15:20.0 | | | | | | | | | | | | | | | | | |
| 9 | 254000000004 | 791 | commodityproduct | Sunflower Seed | 15:20.0 | | | | | | | | | | | | | | | | | |
| 10 | 254000000004 | 825 | procedurecommoditycat | Cereals | 07:58.0 | | | | | | | | | | | | | | | | | |
| 11 | 254000000004 | 825 | procedurecommodity | Maize | 07:58.0 | | | | | | | | | | | | | | | | | |
| 12 | 254000000004 | 825 | proceduredest | UGA->KEN | 07:58.0 | | | | | | | | | | | | | | | | | |
| 13 | 254000000004 | 825 | procedureorigin | EAC | 07:58.0 | | | | | | | | | | | | | | | | | |
| 14 | 254000000004 | 1113 | commoditymarket | Eldoret | 15:45.0 | | | | | | | | | | | | | | | | | |
| 15 | 254000000004 | 1113 | commodityproduct | | | | key_name | sess_id | cell_num_id | age | border | commoditymarket | commoditymetacat | commoditymetacountry | commodityproduct | crossingfreq | education | ... | gender | language | occupation | procedure |
| 16 | 254000000004 | 1339 | commoditymarket | | | | 0 | 50.0 | 2540000000003 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | NaN | NaN | NaN | |
| 17 | 254000000004 | 1339 | commodityproduct | | | | 1 | 52.0 | 2540000000012 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | NaN | NaN | NaN | |
| 18 | 254000000004 | 1347 | procedurecommoditycat | | | | 2 | 54.0 | 2540000000010 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | NaN | NaN | NaN | |
| 19 | 254000000004 | 1347 | procedurecommoditycat | | | | 3 | 59.0 | 2540000000007 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | NaN | NaN | NaN | |
| 20 | 254000000004 | 1347 | procedurecommodity | | | | 4 | 61.0 | 2540000000004 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | NaN | NaN | NaN | |
| 21 | 254000000004 | 1347 | procedurecommodity | | | | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 22 | 254000000004 | 1347 | procedurecommodity | | | | 369359 | 466900.0 | 254000119187 | NaN | NaN | NaN | NaN | NaN | Dry Maize | NaN | NaN | ... | NaN | Swahili | NaN | |
| 23 | 254000000004 | 1347 | procedurecommodity | | | | 369360 | 466901.0 | 254000119188 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | NaN | English | NaN | |
| 24 | 254000000004 | 1347 | procedurecommodity | | | | 369364 | 466906.0 | 254000119190 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | NaN | English | NaN | |
| 25 | 254000000004 | 1347 | procedurecommodity | | | | 369367 | 466909.0 | 254000119191 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | NaN | Luganda | NaN | |
| 26 | 254000000004 | 1347 | procedurecommodity | | | | 369376 | 466921.0 | 254000119192 | NaN | NaN | Kapkwen | NaN | NaN | Rosecoco Beans | NaN | NaN | ... | NaN | English | NaN | |

86080 rows x 21 columns

86080 rows x 21 columns

SUGGESTIONS

| cell_num_id | sess_id | key_name | value_name | created_date |
|--------------|---------|-----------------------|------------|--------------|
| 254000000003 | 50 | procedurecommoditycat | Cereals | 47:17.0 |
| 254000000003 | 50 | procedurecommodity | Maize | 47:17.0 |
| 254000000003 | 50 | proceduredest | UGA->KEN | 47:17.0 |
| 254000000003 | 50 | procedureorigin | EAC | 47:17.0 |
| 254000000012 | 52 | procedurecommoditycat | Cereals | 47:33.0 |
| 254000000012 | 52 | procedurecommodity | Maize | 47:33.0 |
| 254000000010 | 54 | procedurecommoditycat | Cereals | 47:44.0 |
| 254000000010 | 54 | procedurecommodity | Maize | 47:44.0 |
| 254000000007 | 59 | procedurecommoditycat | Cereals | 48:16.0 |
| 254000000007 | 59 | procedurecommodity | Maize | 48:16.0 |
| 254000000007 | 59 | proceduredest | UGA->KEN | 48:16.0 |
| 254000000007 | 59 | procedureorigin | EAC | 48:16.0 |

Demographics

| cell_num_id | border | age | gender | education | crossingfreq | occupation | produce | whatsapp |
|--------------|--------|-------|--------|---------------------|--------------|------------|---------|----------|
| 254000000003 | | 31-40 | Female | University/College | | | | |
| 254000000012 | Busia | 21-30 | Male | Primary | | | Yes | |
| 254000000010 | | 31-40 | Female | No formal education | | Trader | | |
| 254000000007 | Busia | 31-40 | Male | University/College | Daily | Trader | | |

Session Based

| sess_id | cell_num_id | procedurecommoditycat | procedurecommodity | proceduredest | procedureorigin | commoditymarket | commodityproduct |
|---------|--------------|-----------------------|--------------------|---------------|-----------------|-----------------|------------------|
| 50 | 254000000003 | Cereals | Maize | UGA->KEN | EAC | | |
| 52 | 254000000012 | Cereals | Maize | | | | |
| 54 | 254000000010 | Cereals | Maize | | | | |
| 59 | 254000000007 | Cereals | Maize | UGA->KEN | EAC | | |

SUGGESTION EXPLANATION

Session-Based Segmentation

Propose organizing data by `sess_id`, which likely represents a session identifier. We want customer interactions to be tracked in discrete sessions, and have each session associated with specific actions or transactions (e.g., choosing a commodity like 'Maize'). Segmenting customers by sessions can provide insights into customer behavior within individual interactions with the system or service.

Linking Demographic to Sessions

Connect demographic information to the `cell_num_id` to serve as a unique identifier for each customer or transaction. By linking demographics such as age, gender, education, border (which might indicate the crossing point for a transaction), crossingfreq (how often they cross the border), occupation, produce (the type of goods being handled), and usage of whatsapp, a richer profile of each session is established.

07 Enhancing Customer Segmentation: Demographic and Session Data

Predicting the response_theme based on the input features from the cleaned demographic and session data

PREPROCESSING DATA

1. Split Key/Value Pairs

- Pivot tables suggests that the data may need to be organized in a tabular format
- The process involves splitting key-value pairs
- Restructure or normalize the data.

3. Utilize Distinct Entries

- Filtering or processing the data to ensure that each value used for making predictions is distinct and non-repetitive.
- Reducing noise to improve the quality of the predictions
- Focus on the most relevant and unique information available.

2. Redundancy Eliminated

- Merging multiple tables results in duplicate columns that can confuse models or skew results.
- This cleaning process involves removing these redundant columns to streamline the dataset.
- Prevent any potential issues that could arise from having multiple columns with the same data.

PRE-PROCESSING DATA

```
# Step 1: Merge the interaction data
interaction_data = pd.merge(requests, replies, on='request_id')
# Drop the unwanted columns
interaction_data.drop(['update_y', 'sess_id_y'], axis=1, inplace=True)
# Rename the columns
interaction_data.rename(columns={'update_x': 'update', 'sess_id_x': 'sess_id'}, inplace=True)
interaction_data.head()
```

```
interaction_data1 = pd.merge(interaction_data, screen_dict, on='response_id')
interaction_data1.head()
```

```
# Merge the session data
session_data = pd.merge(sessions, interaction_data1, on='sess_id')
# Drop the unwanted columns
session_data.drop(['platform_id_y'], axis=1, inplace=True)
# Rename the columns
session_data.rename(columns={'platform_id_x': 'platform_id'}, inplace=True)
session_data.head()
```

```
# Expanding the databank table
demographics = databank.pivot_table(index=['sess_id', 'cell_num_id'], columns='key_name', values='value_name', aggfunc='first')
demographics.head()
```


PRE-PROCESSING DATA

```
# If you want to drop duplicates based on 'sess_id' and 'cell_num_id' only
demographics_unique = demographics_reset.drop_duplicates(subset=['cell_num_id'])
demographics_unique
```

```
demographic_data = pd.merge(session_data, demographics_unique, on='cell_num_id')
demographic_data
```

```
demographic_filtered = demographic_data[demographic_data['response_theme'] != 'Main Screen']
demographic_filtered
```

```
# Drop the unwanted columns
demographic_filtered.drop(['update_y', 'sess_id_y'], axis=1, inplace=True)
# Rename the columns
demographic_filtered.rename(columns={'update_x': 'update', 'sess_id_x': 'sess_id'}, inplace=True)
```


BUILDING ML PIPELINE

```
# Define your feature DataFrame and target variable
```

```
X = demographic_data.drop(['response_theme', 'reply_id', 'request_id', 'response_id', 'parent', 'created_date', 'notes'], axis=1)
y = demographic_data['response_theme']
```

```
# Split the data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Define categorical and numerical columns
```

```
categorical_columns = X.select_dtypes(include=['object']).columns
numerical_columns = X.select_dtypes(include=['int64', 'float64']).columns
```

```
# Create the preprocessing pipelines for both numerical and categorical data
```

```
numerical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])
```

```
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])
```

RUNNING MODEL

```
# Combine preprocessing steps
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numerical_transformer, numerical_columns),
        ('cat', categorical_transformer, categorical_columns)
    ])

pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', LinearSVC(random_state=42))
])

# Fit the pipeline to the training data
pipeline.fit(X_train, y_train)

# Predict on the test data
y_pred = pipeline.predict(X_test)

# Evaluate the model
print("Classification Report:\n", classification_report(y_test, y_pred))
print("Accuracy Score:", accuracy_score(y_test, y_pred))
```

RESULTS

```
# Evaluate the model
print("Classification Report:\n", classification_report(y_test, y_pred))
print("Accuracy Score:", accuracy_score(y_test, y_pred))
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------------------------|-----------|--------|----------|---------|
| Agriculture | 0.80 | 0.71 | 0.75 | 7219 |
| COVID Info | 0.91 | 1.00 | 0.95 | 18614 |
| Crop Nutrition | 0.62 | 0.20 | 0.31 | 657 |
| Crop/Poultry Farming | 0.77 | 0.35 | 0.48 | 1906 |
| Exchange Rates | 0.89 | 0.83 | 0.86 | 20752 |
| Financial Management | 0.80 | 0.82 | 0.81 | 13075 |
| Good Agricultural Practices | 0.64 | 0.16 | 0.26 | 1523 |
| Health Info | 0.89 | 0.63 | 0.73 | 2621 |
| Health/Legal Info | 0.80 | 0.92 | 0.86 | 4533 |
| Kenya Business Operation | 0.77 | 0.70 | 0.73 | 2849 |
| Kenya Business Operations Info | 0.43 | 0.42 | 0.43 | 48 |
| Legal and Anticorruption Info | 0.75 | 0.63 | 0.68 | 1056 |
| Main Screen | 1.00 | 1.00 | 1.00 | 154129 |
| Market Prices | 0.94 | 1.00 | 0.97 | 86232 |
| Marketplace | 0.97 | 1.00 | 0.98 | 22229 |
| Pests and Diseases | 0.83 | 0.54 | 0.65 | 874 |
| Report Border Experience | 0.80 | 0.60 | 0.69 | 2141 |
| Risk Management/ Insurance | 0.94 | 0.50 | 0.65 | 389 |
| Trade Info | 0.70 | 0.48 | 0.57 | 5737 |
| Weather | 0.91 | 0.91 | 0.91 | 54531 |
| accuracy | | | 0.94 | 401115 |
| macro avg | 0.81 | 0.67 | 0.71 | 401115 |
| weighted avg | 0.94 | 0.94 | 0.94 | 401115 |

Accuracy Score: 0.9418271568004188

CLASSIFICATION RESULTS

- Accuracy: High at 94.18%, indicating strong overall model performance.
- COVID Info & Market Prices: Exceptional precision and recall, almost perfect classification.
- Agriculture: Good performance with a balance between precision (80%) and recall (71%).
- Crop Nutrition: Suboptimal, with low recall (20%), indicating many actual instances were missed.
- Trade Info: Moderate precision (70%) but lower recall (48%), suggesting challenges in identifying all relevant instances.
- Macro Average: Fair performance with equal weighting across categories (Precision: 0.81, Recall: 0.67, F1-Score: 0.71).
- Weighted Average: Excellent, considering class imbalance (Precision: 0.94, Recall: 0.94, F1-Score: 0.94).